

Enabling Remote Direct Memory Access in XRootD

Patrick Adolf, Nazar Burmasov

XRootD overview

Remote Direct Memory Access mechanism

The librdmacm library

OpenFabrics Interfaces

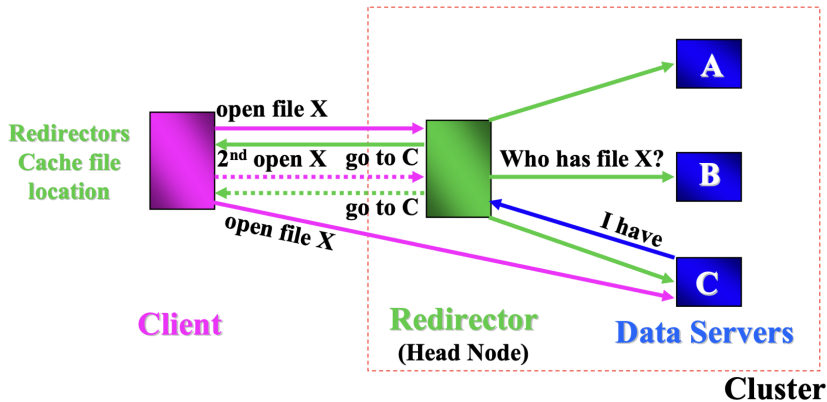
What is XRootD?

- Also called «data access system»
- Fully generic suite for fast, low latency and scalable data access
- High performance, scalable, fault tolerant access to data repositories of many kinds
- Can serve nearly any kind of data:
 - Organized as a hierarchical filesystem-like namespace
 - Based on the concept of directory
- Plugin-based and open
- Dynamic file discovery
 - No record where the files are

What can be done with XRootD?

- One network connection per client for many open files
- Scalable to hundreds of servers
- Allows to reuse files descriptors between clients
- Widely used in the High Energy Physics community
- HEP experiments start using the HPC infrastructure
 - More important for the framework to support Remote Direct Memory Access

Overview

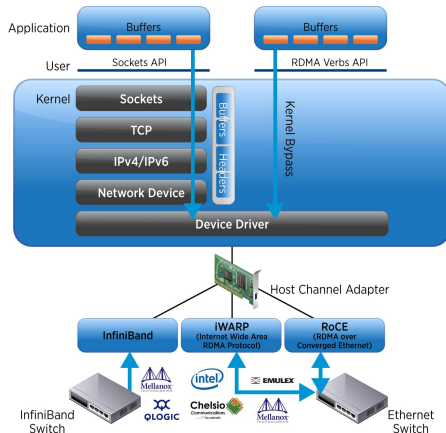


Slurm

- CERN Linux Cluster
- Multi node execution and fast interconnects
- Over 65 000 physical nodes
- Small testing cluster for our project
 - IWARP test setup with two nodes
- Submit jobs to the HPC queue with 'srun' utility
- Process being stuck in D-state
 - One of the servers gets inaccessible

Remote Direct Memory Access

- «**Traditional**» **data transfer**: data must go through buffers using the sockets API → involves TCP, IP stack, device driver → CPU overhead
- **RDMA**
 - OS interventions and buffers are avoided
 - Zero-copy networking – no work done by CPUs
 - High throughput, low-latency communications → good for high-performance computing
 - Needs dedicated Host Channel Adapters



librdmacm

- Library for managing RDMA connections
 - Allows applications to set up reliable connected and unreliable datagram transfers over RDMA
- Transport-neutral interface
 - Same code can be used for both InfiniBand and IWARP adapters
 - Huge benefit
- Interface based on sockets; adapted for queue pair based semantics
- Communication must use a specific RDMA device
- Data transfers are message-based

librdmacm

- **librdmacm** only provides communication management
 - Connection setup and tear-down
- **librdmacm** is build on top of **libverbs**
 - Responsible for carrying out the transfer

How to send data with librdmacm

- Focus on a minimalistic example code to transfer a buffer
 - Base for further developments
- [Link: code on Github](#)
- Uses build.sh script to build the project
- Usage:
 - Server: `srun -w [hostname] -p batch-short -t 00:05:00 -n 1 /path/to/our/code/rcopy`
 - Client: `srun -w [hostname] -p batch-short -t 00:05:00 -n 1 /path/to/our/code/rcopy [servername]`
- [Link: rcopy.c code](#)

rcopy.c

- Main function is splitted into a server and a client code
- `server_opt` and `client_opt` are only for parsing command line arguments

```
1  int main(int argc, char **argv){
2      int ret;
3      if (argc == 1 || argv[1][0] == '-') {
4          server_opts(argc, argv);
5          ret = server_run();
6      } else {
7          client_opts(argc, argv);
8          ret = client_run();
9      }
10     return ret;
11 }
```

Server Function

```
1  static int server_run(void)
2  {
3      int lrs, rs;
4      union rsocket_address rsa;
5      socklen_t len;
6
7      lrs = server_listen();    // checking server address
8      if (lrs < 0)
9          return lrs;
10     len = sizeof rsa;
11     printf("waiting for connection...");
12     fflush(NULL);
13     rs = raccept(lrs, &rsa.sa, &len);    // connecting server
14     printf("client: %s\n", _ntop(&rsa));
15
16     char buffer[17];           // buffer for the recieving data from client
17     memset( buffer, 0, sizeof(buffer) );
18     _recv(rs, buffer, sizeof(buffer)-1 ); //Recieving the sending data from client
19     buffer[16] = 0;
20     printf("received buffer: %s", buffer);
21
22     char buffer2[16];          // second buffer for own data sending to client
23     memset( buffer2, 'B', sizeof( buffer2 ) );
24     len = rsend(rs, buffer2, sizeof( buffer2 ), 0); //Sending data with information about size and socket tu client
25     printf( "send %d bytes", len );
26
27     rshutdown(rs, SHUT_RDWR);
28     rclose(rs); //end connection
29     return 0;
30 }
```

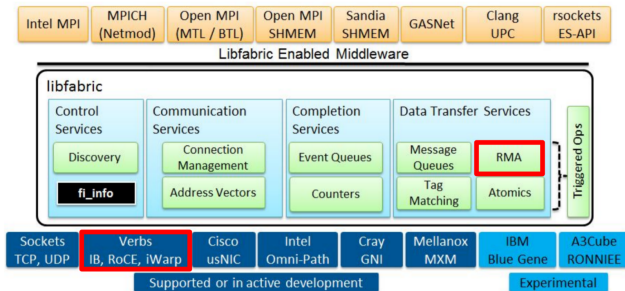
Client Function

```
1  static int client_run(void)
2  {
3      struct msg_hdr ack; // declaring some needed variables
4      int ret, rs;
5      size_t len;
6
7      rs = client_connect(); // connecting client
8      if (rs < 0)
9          return rs;
10     printf("...");
11     fflush(NULL);
12     gettimeofday(&start, NULL);
13
14     char buffer[16]; // buffer with own data to send to server
15     memset( buffer, 'A', sizeof( buffer ) );
16     len = rsend(rs, buffer, sizeof( buffer ), 0); //Sending own data with information about socket and size
17     printf( "send %d bytes", len );
18
19     char buffer2[17]; // empty buffer for the data recieving from server
20     memset( buffer2, 0, sizeof(buffer2) );
21     _recv(rs, buffer2, sizeof(buffer2)-1 ); //Getting data from server
22     buffer2[16] = 0;
23     printf("received buffer: %s", buffer2);
24     gettimeofday(&end, NULL);
25
26     shutdown:
27     rshutdown(rs, SHUT_RDWR);
28     rclose(rs); // end connection
29     return 0;
30 }
```

OpenFabrics Interfaces

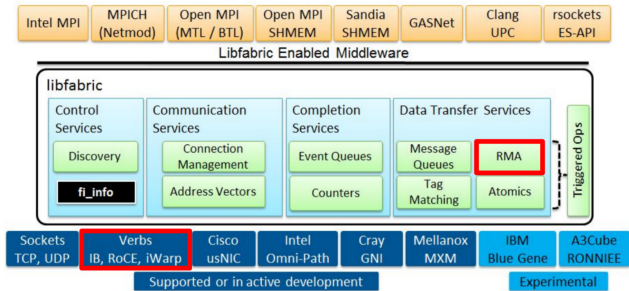
The **libfabric** library

- Provides interfaces to a number of hardware solutions and some basic services
- In most cases, applications access the provider implementation directly → **low latency**
- Providers: hardware specific; plug into the framework to provide access to the fabric hardware



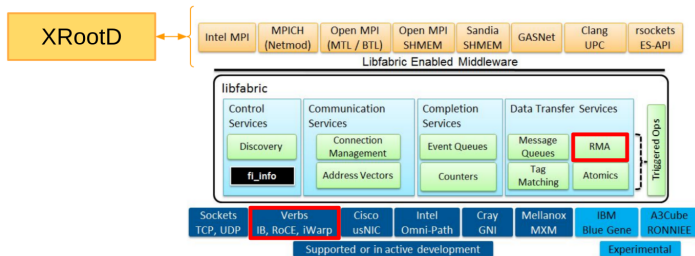
Main parts of the libfabric:

- **Control services:** discover system communications
- **Communication services:** setup communication between nodes
- **Completion services:** report results of operations (event queues or counters)
- **Data transfer services:** set of interfaces for different communication approaches



Enabling XRootD:

- Need to implement a simple example for further developments



RDMA file copying with libfabric

- Simple file transfer using [libfabric](#) and some modified utilities from [fabtests](#)
 - Github repository: [libfabric](#)
 - Github repository: [fabtests](#)
- Built with CMake to enable automatic fetching of in-system [libfabric](#) version
- Simple usage:

Client: `./oficp -r[ead] /full/path/to/file -y[ank] /full/path/to/copy hostname`

Server: `./oficp`

- The full code is available on Github

Client code

- `ft_get_tx_comp()` is implemented in terms of `fi_cntr_read()` and `fi_cntr_wait()`, which allows interacting with the completion queue

```
1 FILE* file_ptr = fopen(opts.src_filename, "rb");           // open the file for reading
2 fseek(file_ptr, 0, SEEK_END);                               // get file's size
3 size_t file_len = ftell(file_ptr);
4 snprintf(tx_buf, tx_size, "%s", opts.dst_filename);
5 int name_len = strlen(opts.dst_filename) + 1;
6 (void) fi_write(ep, tx_buf, name_len, mr_desc,              // send destination
7               remote-fi_addr, remote.addr, remote.key,      // file name to the server
8               &fi_ctx_write);
9 (void) ft_get_tx_comp(++tx_seq);
10 char file_len_ch[256];
11 snprintf(file_len_ch, sizeof(file_len_ch), "%zu", file_len);
12 snprintf(tx_buf, tx_size, "%s", file_len_ch);
13 fi_write(ep, tx_buf, sizeof(file_len_ch), mr_desc,          // send file size to the server
14         remote-fi_addr, remote.addr, remote.key,
15         &fi_ctx_write);
16 (void) ft_get_tx_comp(++tx_seq);
17 rewind(file_ptr);
18 tx_buf = (char*)malloc((chunk_size + 1) * sizeof(char)); // allocating chunk
19 size_t n_transx = file_len / chunk_size;                  // number of transactions
20 if (file_len % chunk_size != 0)
21     n_transx++;
22 for (size_t itx = 0; itx < n_transx; itx++) {               // read and send the file chunk by chunk
23     for (size_t i = 0; i < chunk_size; i++)
24         tx_buf[i] = '\0';
25     fread(tx_buf, chunk_size, 1, file_ptr);
26     (void) fi_write(ep, tx_buf, chunk_size, mr_desc,
27                     remote-fi_addr, remote.addr, remote.key,
28                     &fi_ctx_write);
29     (void) ft_get_tx_comp(tx_seq++);
30 }
31 fclose(file_ptr);                                           // close the file and wait
32 (void) ft_get_rx_comp(rx_seq++);                             // for the success message
33 fprintf(stdout, "Received data from Server: %s\n", (char*)rx_buf);
```

Server code

- `ft_get_rx_comp()` is also implemented in terms of `fi_cntr_read()` and `fi_cntr_wait()`

```
1  (void) ft_get_rx_comp(rx_seq++);           // receive destination file name and size
2  opts.dst_filename = (char*)rx_buf;        // and prepare to write
3  FILE* dst_file_ptr = fopen(opts.dst_filename, "wb");
4  size_t file_len;
5  for (size_t i = 0; i < rx_size; i++)
6      rx_buf[i] = '\0';
7  (void) ft_get_rx_comp(rx_seq++);
8  file_len = atol(rx_buf);
9  size_t n_transx = file_len / chunk_size;
10 if (file_len % chunk_size != 0)
11     n_transx++;
12 for (size_t irx = 0; irx < n_transx; irx++) { // receive and write the file
13     for (size_t i = 0; i < chunk_size; i++) // chunk by chunk
14         rx_buf[i] = '\0';
15     (void) ft_get_rx_comp(rx_seq++);
16     fwrite((char*)rx_buf, chunk_size, 1, dst_file_ptr);
17 }
18 fclose(dst_file_ptr);                       // close the file and send
19 char* fin_message = "Success: end of transmission"; // the completion message
20 size_t msg_size = strlen(fin_message);
21 snprintf(tx_buf, tx_size, "%s", fin_message)
22 (void) fi_write(ep, tx_buf, msg_size, mr_desc,
23               remote-fi_addr, remote.addr, remote.key,
24               &fi_ctx_write);
25 (void) ft_get_tx_comp(tx_seq++);
```

Testing file copying

- Test data: English Wikipedia dump (2006)
 - enwik8** – $100 \cdot 10^6$ bytes
 - enwik9** – $1000 \cdot 10^6$ bytes
- The chunk size per transfer has to be defined manually – using 10^6 bytes for this test

Client

```
nburmaso hpc003 /hpcscratch/user/nburmaso/project/oficp/build
→ ./oficp -r /hpcscratch/user/nburmaso/project/oficp/data/text_samples/enwik8 -y /hpcscratch/user/nburmaso/project/oficp/data/enwik8_copy hpc002
Trying to read a file: /hpcscratch/user/nburmaso/project/oficp/data/text_samples/enwik8
file_len = 100000000
RMA write to server, chunk #0
RMA write to server, chunk #1
RMA write to server, chunk #2
(...)
RMA write to server, chunk #99
Received data from Server: Success: end of transmission
```

Server

```
nburmaso hpc002 /hpcscratch/user/nburmaso/project/oficp/build
→ ./oficp
Received data from Client: /hpcscratch/user/nburmaso/project/oficp/data/enwik8_copy
Opening file for writing: /hpcscratch/user/nburmaso/project/oficp/data/enwik8_copy
file_len = 100000000
Received data from Client: #0
Received data from Client: #1
Received data from Client: #2
(...)
Received data from Client: #99
```

Conclusion

- [RDMA](#) is a powerful mechanism, which can be used to implement a high-throughput, low-latency infrastructure
- Adopting [librdmacm](#) for [XRootD](#) is a low-hanging fruit
- A simple file transfer is implemented using [libfabric](#), and it can be used for further developments